

BUI THE TRUNG

Cau Giay, Hanoi, Vietnam

(+84) 373 104 304 | trune.naee@gmail.com | <https://github.com/Trungnef>

BACKGROUND

AI Engineer focused on LLM fine-tuning, Retrieval-Augmented Generation (RAG), AI agents, and Vietnamese NLP. Experienced in building production-oriented AI applications for document QA, administrative text generation, semantic search, and conversational workflows. Research interests include reasoning in Small Language Models, test-time scaling, and efficient NLP systems.

EDUCATION

Vietnam Japan University (VJU) - Vietnam National University, Hanoi 2021 - 2025

Bachelor of Computer Science and Engineering

GPA: 3.61/4

Relevant coursework: Algorithms for Data Science; Natural Language Processing; Machine Learning; Deep Learning; Statistical Methods for Data Analysis.

WORK EXPERIENCE

Aimesoft Company, Hanoi - AI Engineer Dec 2025 - Present

NTTEA - Japanese Administrative Text Generation

- Built data preparation and fine-tuning pipelines for Japanese administrative writing tasks.
- Fine-tuned Tsuzumi-28B with LoRA on datasets derived from 120+ Japanese government regulations.
- Developed rewriting and document-generation workflows to produce compliant administrative documents from structured inputs.

Tech stack: Python, PyTorch, Hugging Face Transformers, PEFT/LoRA, LlamaFactory, Tsuzumi-28B, Azure, Dify.

Rakuyaku - QC Check / Consistency Check Tool

- Built AI backend services for consistency checking across multi-document PDF collections.
- Implemented a RAG pipeline with PDF parsing, hybrid retrieval, query fusion, and context filtering to support semantic consistency checks.
- Integrated vector storage, caching, and LLM-based response generation for document QA and review workflows.

Tech stack: Python, FastAPI, LlamaIndex, LlamaParse, Azure OpenAI, Qdrant, Redis, BM25/vector hybrid search, Docker.

RAG (QA Assist)

- Developed an LLM orchestration service for document-based, multi-turn question answering.
- Implemented document reference management, file caching, and chat-history formatting to preserve context across conversations.
- Optimized file handling by reusing valid AI file references and processing long-running upload/API operations asynchronously.

Tech stack: Python, FastAPI, Google Gemini API, GenAI File API, S3, async processing.

Linebot Meisei - LINE Mini App Ordering Bot

- Developed backend services and product features for a LINE-based ordering bot integrated with a LINE Mini App.
- Built LangChain-based agents with tool-calling for product inquiry and order-placement workflows.

Tech stack: Python, FastAPI, LangChain, tool-calling agents, LINE Messaging API, LINE Mini App, Docker.

DataCom Corp, Japan - AI Engineer Mar 2025 - Jun 2025

Retail Manufacturer Matching Platform

- Developed a platform connecting manufacturers, brands, and retailers in the CPG industry.
- Built data-ingestion pipelines for Amazon product information and implemented LLM-based matching and recommendation logic.

Tech stack: Python, data ingestion, Amazon product data processing, embeddings, semantic matching, LLM-based recommendation.

RESEARCH EXPERIENCE

Research on Chain-of-Thought Techniques in Large Language Models <i>Vietnam Japan University - Hanoi Student Research</i> - Compared zero-shot, few-shot, CoT, self-consistency, and ReAct prompting on Vietnamese reasoning problems using local and hosted LLMs. - Implemented accuracy and inference assessments to analyze trade-offs between reasoning quality, cost, and latency.	Feb 2025 - Jun 2025
Vietnamese Language Model with Mamba-2-Hybrid Architecture <i>Vietnam Japan University - Hanoi Student Research</i> - Developed a Vietnamese language model using the Mamba-Hybrid architecture and extended Falcon Mamba 7B for Vietnamese vocabulary and syntax. - Evaluated long-context performance and presented research findings at academic and professional forums.	Jul 2024 - Nov 2024

PUBLICATIONS

Bridging the Reasoning Gap in Vietnamese with Small Language Models via Test-Time Scaling - STDJ <i>Authors: Bui The Trung, Do Minh Duc, Nguyen Van Vinh, Bui Nguyen Quoc Trinh</i> - Built Vi-S1K for training and Vi-Elementary-Bench for evaluating Vietnamese mathematical reasoning. - Showed that supervised fine-tuning improved explanation quality by 77% on a 1.7B-parameter model; analyzed CoT, self-consistency, and ReAct prompting trade-offs.	2026
---	------

SKILLS

Programming	Python
LLM & AI	RAG, LLM fine-tuning (LoRA), prompt engineering, AI agents, semantic search, document QA, NLP, machine learning, deep learning
Frameworks & Tools	PyTorch, Hugging Face, LangChain, LlamaIndex, LlamaFactory, Dify, Git, Docker, FastAPI.
Cloud Platforms	Azure, AWS, GCP

SELECTED HONORS & ACTIVITIES

Encouraging Studying Scholarship	2021 - 2025
Outstanding Young Face of VJU	2023
Toshiba Scholarship	2023
BIDV Scholarship	2021
10th Youth Global Leadership Program, Seoul, Korea	2024
Certificate of Commendation, Youth Union of Vietnam National University / VJU	2024
Second Place, Law Propagation Contest "Gender Equality Through the Lens of Students"	2023

CERTIFICATIONS

IELTS - 6.5	2025
-------------	------